

CONCLUSION AND SUMMARY

This Chapter highlights the contributions, limitations and personal insights of the author in this thesis and proposes possible directions for future research.

6.1 Main Contributions

This thesis specifically addressed decision making by agents under uncertainty. In the field of Robotics and Artificial Intelligence (AI), considering uncertainty in search policies is not straightforward. This is due to the complexity involved in solving Partially Observable Markov Decision Processes (POMDP), commonly used to describe uncertainty in tasks, which become quickly infeasible for even the simplest problems. Although there has been progress in the development of POMDP solvers with demonstrated applications to robotics, these are predominantly verified in simulation where the action space is discretised.

This thesis concentrates on the learning of search policies from human teachers and their transfer to a robotic apprentice, and provides three major contributions to the current research which will be summarised below.

Firstly, we demonstrate a Programming by Demonstration POMDP framework (PbD-POMDP) to learn a mixture of search strategies from human teachers. We use a particle filter to represent the belief of the end-effector's position of both human teachers and the robot apprentice. The particle filter is compressed to a belief state composed of the most likely state and entropy. A sequence of demonstrated belief states and actions is used to learn a generative joint belief-action distribution parameterised by Gaussian Mixture Model (GMM). The GMM is used as an autonomous dynamical system which gives a velocity vector field when conditioned on the current belief state. We demonstrate that mixed behaviours, such as risk prone and adverse, are encoded by the generative model.

Secondly, we propose a Reinforcement Learning (RL) approach to further optimise the parameters of the GMM model to take into account the quality of the demonstrated behaviours. By introducing a binary cost function and using approximate dynamic programming in an actor-critic framework (which we call RL-PbD-POMDP) we are able to significantly improve the performance of a search strategy designed to accomplish a peg-in-hole task.

Thirdly, we develop a Bayesian State Space Estimator (BSSE), which we name Measurement Likelihood Memory Filter (MLMF), to solve the Simultaneous Localisation and Mapping (SLAM) problem where only haptic information is available. Given the nature of the searches, we consider any prior structure in the beliefs (such as Gaussianity) to be ill-suited. We demonstrate that by choosing a different parameterisation of the Histogram Bayesian filter we are able to overcome the inherent exponential space and time complexity. To the best of our knowledge this is the first work which considers mostly negative information in a SLAM setting.

6.2 Limitations and Future Work

We summarise the current perceived limitations of our work and discuss different approaches to resolving them.

GAUSSIAN MIXTURE CONTROLLER

Throughout this dissertation, we used Gaussian Mixture Models (GMM) to encode a vector field policy which is a function of the current belief state. For the robot to be able to localise itself the policy has to be guided towards salient areas of the state space, such as edges and corners. However, as the GMM learning is non-deterministic and data driven, there are no guarantees that behaviour such as remaining in contact with edges in order to get haptic feedback will be encoded and reproduced by the GMM policy.




Figure 6.1: GMM policy drawback. *Top:* Two demonstrations, blue trajectories, follow a path which when fitted with a GMM (one Gaussian component is drawn in green) results in a vector which keeps the plug in contact with the edge of the wall. *Bottom:* A third demonstration, red line, did not result in a contact with the edge of the wall. The Gaussian of the GMM is shifted to the mean point of the data. The GMM policy no longer keeps the plug in contact with the edge.

It would require only one poorly demonstrated trajectory which failed to

establish a contact with an edge, for the vector field to be shifted resulting in a policy which fails to establish contact with the environment. Figure 6.1 illustrates an example of this drawback.

A solution would be to design a cost function which gives more importance to data points which are close to salient features, such as edges and corners. As a result, the components of the problematic red trajectory in Figure 6.1 are excluded during the learning. This can be integrated into the Actor-Critic Reinforcement Learning (RL) framework used in Chapter 4, for the peg-in-hole task.

RL is an off-line approach (in the sense that many rollouts are needed to achieve the desired behaviour) which can be used to select behaviour which will result in a policy remaining in contact with the environment. However, if the geometry of the environment was to change by less than a centimetre the same situation would occur (failure to remain in contact) and GMM policy parameters would have to be relearned through either new demonstrations or via autonomous exploration.

In the peg-in-hole task described in Chapter 4, in order to enforce a constant contact with the environment, a hybrid force/position controller was used which disregarded the velocity component orthonormal to the wall. The remaining two velocity components were obtained from the GMM policy and modulated by a heuristic function to surmount the edges of the plug.

The main problem arises however during the execution of the GMM policy when no sensory feedback constraints are resolved. Belief space planners (reviewed in Chapter 2) are approaches which take into consideration variations in the environment and which can produce trajectories which actively seek sensory feedback. These planners solve an objective function online, with contact constraints [Vien and Toussaint \(2015b\)](#). Such online belief space planners are computationally expensive and require simulation of dynamics and sensory feedback in order to find an appropriate set of actions. A dynamic system approach, such as the GMM policy, learns the sensory-control mapping directly making it computationally cheap at runtime.

Future research could consider local adaptation of the dynamical system in order to try and seek out specific sensory feedback. This could be achieved for instance by combining planning with GMM policies or learning local dynamical systems which seek out sensory feedback. The difficulty lies in combining their joint actions.

BELIEF COMPRESSION

Throughout this research, the belief is represented by a probability density function and the GMM policy is learned from a dataset with a fixed number of dimensions, which is seven in our case (3 for velocity, 3 for the most likely state and 1 for entropy). The compression of the belief to a belief state vector

results in loss of information which weakens the Markov assumption. There exist however other compression methods, such as E-PCA [Roy and Gordon \(2003b\)](#). This is a non-linear dimensionality reduction technique which retrieves a set of basis probability distributions in order to minimise the reconstruction Kullback–Leibler (KL) divergence.

However, such compression methods require a discretisation of the probability density function and then its projection to a latent space, which in the case of E-PCA, requires solving an convex optimisation problem (iterative Newton’s method) at each time step.

Furthermore, it is also not clear what effect an improved belief compression method would have on the policy. The better the compression method (in terms of KL divergence) the more dimensions are necessary and as a result more data points will be required to train the GMM. We make the observation in both Chapter 3 and Chapter 4 that the GMM policy is better than a myopic policy. However, this holds true only when a large amount of uncertainty is present. When the uncertainty starts to decrease the Greedy method performs just as well or even better than a four dimensional belief state GMM. As such, it is not clear that a more sophisticated belief compression method for the tasks we consider would be an improvement.

POLICY REPRESENTATION

We learned the belief space policies in task space (Cartesian position in the world’s frame of reference). This choice entails two difficulties: the number of parameters needed to encode the task and generalisation.

As there is much variance in the demonstrated search behaviour (at the raw trajectory level) many parameters are necessary to encode the policy. The policies learned for the search tasks in this thesis required around 90 Gaussian functions of dimension 7. This is a considerable number of parameters considering the simplicity of the task (find the the table/wall, navigate to an edge, reach the goal). Typically more parameters also entail a poor generalisation.

Future work could be directed towards learning a high-level policy composed of parameterised behaviour policies which are easily re-usable. Such policies could be parameterised by sub-goals and contact constraints which could be extracted by segmenting the original demonstrations.

MLMF

The Measurement Likelihood Memory Filter (MLMF), introduced in Chapter 5, is based on the assumption of a sparse likelihood function where mostly zero measurements are present. This by itself is not a weakness, but a necessary assumption to achieve a linear space complexity. As the time complexity remains exponential with respect to the number of objects (although at a lesser growth

rate) we were obliged to introduce an additional independence assumption. This assumption implies that some information will be lost and the filtered marginals will be an approximate solution with respect to an optimal Bayesian filter. There is no obvious remedy to this problem. There are two approaches: either introduce an assumption (as we did) or perform an approximate evaluation of the marginalisation. An approximate marginalisation could be achieved by evaluating the value of the marginals at specific points and setting the remaining marginal values by interpolation.

6.3 Final Words

During my PhD I have spent a considerable amount of time studying the role of uncertainty in Artificial Intelligence, specifically how it affects decision making in agents. I list below some points of interest and important insights.

- *Humans can be poor teachers:* Programming by Demonstration (PbD), traditionally requires an expert teacher to provide a set of demonstrations in the form of state-velocity pairs which are used to learn a policy represented by a regressor function. If the teacher is rarely successful at his task, learning a policy directly from his behaviour will yield a policy on par with the teachers performance. One of the original questions posed by PbD is *what to imitate ?*, [Billard et al. \(2008a\)](#). By introducing a simple binary cost function, as shown in Chapter 4, we were able to improve the quality of the policy. All regression based PbD methods can use the Reinforcement Learning (RL) approach used in Chapter 4, with no additional rollouts being necessary.
- *Reinforcement Learning can be easy (continuous state and action space):* RL is notable for needing lengthy simulations and episodes to be successful, which typically result in a complete exploration of the entire state (or parameter) space for tasks such as the inverted-pendulum or mountain cart. This is infeasible for learning a complicated continuous state and action space policy with a long time horizon.

In using RL during both my PhD and Master Thesis, there was always some magic required to get RL to work, such as choosing an appropriate learning rate of the value function and decay rate of the exploration noise, in order to get past local minimas.

There are two factors which make RL difficult to use: the exploration-exploitation dilemma and the non-stationarity of the value function approximator during on-line learning. To alleviate the exploration-exploitation problem one can use human demonstrations in which an optimal set of state-action pairs hopefully exists. The non-stationarity of the target value function can be achieved through batch methods, also known as fitted RL

methods, which keep all data witnessed during episodes and learn the value function off-line through approximate dynamic programming. Learning the value function online at each time step can lead to divergence if an appropriate function approximator is not chosen (Szepesvári, 2010, p. 51). Given the current memory capacity of modern computers the fitted RL off-line methods seem more appropriate since the RL problem becomes a familiar regression problem for which many algorithms are applicable.

By using a fitted off-line approach to learn a value function in combination with a separate parameterisation of the policy in an Actor-Critic framework, it is again feasible to use simple reward functions which can result in significant improvements in the policy, as shown in Chapter 4.

- *Myopic vs POMDP policies:* Across most of the discrete POMDP literature the Q-MDP myopic policy is used as a benchmark. In the problems considered, Q-MDP does not fare as well in terms of reward but the difference in the policies's parameters are not reported. Throughout this research we always used a myopic policy as a benchmark to compare with our more sophisticated approaches. Even though myopic policies do not consider uncertainty when taking an action, the action will always lead to some information gain, since actuated robots are very precise. As a result the uncertainty will always decrease and eventually the robot will manage to accomplish its task. It is not clear what difference exists between myopic and continuous policy search methods on point motion and grasping tasks in which the uncertainty is considered uni-modal and Gaussian, as there has been no extensive comparison. This being said a lot depends on the perception capabilities of the robot. In an environment which provides enough sensory cues after a few time steps the robot will be localised regardless of the policy followed.

